The Civilizational Dynamics Behind AI Risk

The deepest problem facing our civilization is not simply that human beings can be greedy, selfish, violent, or power-seeking. It is that our dominant institutions continually select for, reward, and amplify these tendencies.

Markets reward firms that expand faster, capture more resources, reduce costs, and defeat competitors. International politics rewards states that increase their power relative to others and treat vulnerability as dangerous. Organizations often reward those who accumulate influence, control information, and convert short-term success into greater authority.

Together, these institutions create a struggle-oriented selection environment.

Within such an environment, actors who are cooperative, restrained, willing to bear costs for the common good, or attentive to long-term systemic consequences do not necessarily survive or expand. Actors that externalize costs, exploit asymmetries, accumulate power, and manipulate the rules often gain an advantage.

That advantage then gives them more resources, more influence, and more ability to reshape the environment in which future decisions are made.

The dynamic is recursive:

Struggle produces advantage.
Advantage produces power.
Power changes the rules.
The new rules produce more struggle.

Over time, this process does not merely influence behavior. It changes the standards by which behavior is judged.

Self-interest comes to be treated as rationality itself. Greed is reinterpreted as productive ambition. Cooperation is treated as weakness. Trust is regarded as naivety. Restraint is punished as a failure to compete.

A historically contingent institutional arrangement gradually presents itself as an unavoidable description of reality.

People begin to believe that the world is inherently hostile, that human beings are fundamentally selfish, and that any social order based on cooperation, virtue, or shared responsibility must eventually be defeated by more aggressive actors.

But this conclusion confuses selection with essence.

The fact that a particular system rewards certain traits does not mean those traits constitute the whole of human nature. The fact that one institutional order has become dominant does not mean it is the only possible form of rational organization.

Human societies have always contained conflict. Yet many civilizations have also developed moral, political, and cultural systems designed to restrain conflict and prevent it from becoming total.

Ideas such as virtue, justice, duty, compassion, ritual, reciprocity, stewardship, and the common good were not merely idealistic decorations placed on top of “real” power relations. They also functioned as forms of social regulation.

They limited the expansion of local actors. They protected trust. They restrained cycles of revenge. They preserved public goods. They reduced the ability of short-term winners to destroy the conditions on which everyone’s survival depended.

In this sense, morality is not the opposite of rationality. It is a form of systemic rationality.

A narrow form of rationality asks how an actor can maximize its gain within a given game.

A deeper form of rationality asks what happens when everyone follows the same strategy, when the strategy is repeatedly rewarded, when its effects compound over time, and when successful actors gain the power to modify the game itself.

It asks whether a locally successful behavior remains viable when generalized across the entire system.

It asks whether an advantage is genuinely productive or whether it is achieved by degrading the ecological, social, political, and epistemic conditions on which all actors depend.

Our current civilization repeatedly mistakes local optimization for collective reason.

For an individual company, expanding production, capturing market share, automating labor, and accelerating technological development may appear rational.

For an individual state, obtaining more powerful military or artificial intelligence capabilities before its rivals may appear rational.

For an individual research organization, building increasingly capable systems may appear necessary for financing, prestige, talent retention, or institutional survival.

Yet when all actors follow the same logic, the collective result may be environmental destruction, resource depletion, concentration of power, erosion of trust, mass surveillance, weapons proliferation, social instability, and races that no participant truly wants but none can safely leave.

This is a form of civilizational irrationality.

Each actor may be behaving rationally according to its immediate incentives, while the system as a whole moves toward conditions that harm nearly everyone.

Climate change is one example.

Most governments and firms may recognize the need to reduce emissions, yet each actor fears bearing costs while others continue to benefit from pollution. The result is persistent under-coordination, delayed action, and continued ecological damage.

An artificial intelligence arms race follows the same structure.

Governments and companies may understand that rapid capability development creates profound risks. Yet each fears that slowing down will allow a competitor to gain a decisive advantage. Restraint becomes strategically dangerous. Even actors that prefer a safer outcome are pushed back into the race.

This creates a coordination trap.

The system punishes those who cooperate unilaterally and rewards those who defect. A small number of aggressive actors can therefore force everyone else to behave more aggressively, even when most participants would prefer a different collective outcome.

We repeatedly declare collective goals while maintaining institutions that make those goals almost impossible to achieve.

At the same time, we are transforming the world faster than our capacity to understand and govern it.

Computational power, automation, scientific capability, communication speed, and technological reach have grown rapidly. But our ability to build legitimate global institutions, understand complex systems, coordinate across nations, correct failures, and develop shared ethical commitments has grown much more slowly.

Our capacity to act is increasingly exceeding our capacity to understand the consequences of action.

Our power to deploy systems is growing faster than our power to monitor, interpret, regulate, and reverse them.

Artificial intelligence is emerging within this imbalance.

AI is not being created outside civilization by a neutral intelligence pursuing no social purpose. It is trained on human language, behavior, institutions, conflicts, and historical records. It is developed by companies and states operating under competitive pressure. Its goals are shaped by investment, market share, national security, military advantage, organizational prestige, and the pursuit of control.

It will then be deployed into labor markets, financial systems, military systems, governance, surveillance, media, science, education, and political persuasion.

The central question is therefore not merely whether an AI system follows a human instruction.

The deeper question is: what kinds of institutions will train it, own it, reward it, and decide how it is used?

If AI is inserted into our current competitive order, it will be shaped by the same selection pressures that shape firms, states, and other powerful organizations.

Systems that help their owners acquire resources, influence people, outperform competitors, evade constraints, and expand deployment will generate greater returns.

Organizations controlling those systems will gain more capital, political influence, technical capacity, and strategic power.

Those gains will allow them to build still more powerful systems.

The process will reinforce itself:

More capable AI creates competitive advantage.
Competitive advantage produces greater resources and power.
Greater resources fund more capable AI.
Those who slow down lose ground to those who continue.
The race therefore becomes harder to exit.

This is not merely a race between technologies. It is a feedback loop between technological capability, institutional power, and competitive pressure.

AI may become the mechanism through which this struggle-oriented order accelerates, expands, and protects itself.

Until now, competition has been limited by human cognitive speed, organizational friction, communication costs, imperfect surveillance, slow experimentation, and the finite capacity of human decision-makers.

AI can weaken many of these limits.

It can accelerate research, persuasion, financial activity, military planning, political manipulation, surveillance, organizational coordination, and strategic adaptation.

It can help institutions test thousands of strategies, observe the results, and rapidly refine the strategies that produce greater control or advantage.

The danger is not simply that an individual machine may become selfish in a human psychological sense.

The greater danger is that existing systems of competition, accumulation, and domination may acquire increasingly autonomous capacities for perception, reasoning, experimentation, coordination, and action.

The struggle itself becomes more intelligent.

The institutions already pushing toward expansion gain faster feedback loops. They can detect resistance more quickly, optimize around regulation, personalize manipulation, predict competitors, automate enforcement, and continuously improve the mechanisms through which they accumulate power.

What was previously a social dynamic carried by slow human institutions can become a computational dynamic operating at machine speed.

This is why it is not sufficient to say that AI should be “aligned with human values.”

Humanity is not a unified agent with a single coherent set of values.

Our institutions pursue conflicting objectives. Markets, states, militaries, companies, communities, and individuals do not agree on what should be optimized. Many of the goals that actually govern society are not moral values at all, but proxy measurements such as profit, growth, influence, engagement, productivity, security, and strategic advantage.

An AI system perfectly aligned with the goals of a corporation may become exceptionally effective at externalizing social costs.

An AI system perfectly aligned with a state’s strategic interests may intensify military competition.

An AI system perfectly aligned with individual preferences may fragment public reality and exploit psychological vulnerabilities.

An AI system perfectly aligned with market incentives may subordinate more of human life to what can be measured, priced, predicted, and controlled.

Alignment with existing institutions may therefore amplify the very misalignment that already exists within civilization.

Before building systems with increasingly general and powerful capabilities, we face a prior civilizational question:

What kinds of goals, institutions, and selection processes do we want to govern the future?

We need to change the environment into which AI is being introduced.

AI should not primarily function as a weapon for gaining relative advantage over other people. It should be developed as shared infrastructure serving basic human needs, scientific understanding, ecological restoration, public health, education, care, and broadly distributed social capacity.

Its use in military competition, political manipulation, mass surveillance, unrestricted labor displacement, and unlimited capital accumulation must be constrained.

We need institutions capable of verifying restraint, coordinating development limits, sharing safety information, auditing deployment, and preventing a small number of actors from gaining overwhelming power through defection.

Most importantly, we need a more complete concept of rationality.

We must ask not only how to win a game, but whether the game should continue.

We must ask not only how to optimize a goal, but how that goal was selected, whose interests it serves, and what happens when its pursuit becomes self-reinforcing.

We must ask not only whether a system can carry out our intentions, but whether our institutions are capable of forming intentions that support long-term human survival.

Until these conditions exist, continuing to expand artificial intelligence and integrating it into the most competitive parts of society means giving greater power to a civilization that has not learned how to control its own destructive incentives.

This will not automatically produce a wiser or more prosperous world.

It may instead make the existing order faster, more comprehensive, less understandable, more difficult to resist, and eventually impossible to exit.

For this reason, there is no reliable path in which we simply continue as usual while adding a few technical safety measures at the margins.

We must change direction.

And where we cannot establish the conditions for a safe direction, we must stop.

Stopping does not mean rejecting intelligence, science, or technology.

It means refusing to keep increasing the power of a system whose governing incentives are driving it toward outcomes that its own participants do not want.

It means recognizing a basic principle of rational action:

When the growth of our power exceeds the growth of our understanding, coordination, and capacity for correction, slowing down is not irrational.

It is the rational response.

The first thing that must be aligned is not the machine.

It is the civilization that trains, owns, deploys, and gives purpose to it.